

Springboard Data Science Capstone Project Airbnb Yearly Revenue Prediction

Ashley Jiangyang

Mar, 2020

Table of Contents

1.	<i>Introduction</i>	<i>3</i>
2.	<i>Data Collection and Wrangling Summary</i>	<i>3</i>
2.1	Drop Features	4
2.2	Re-code the Features.....	4
2.3	Deal with Missingness	4
2.4	Outliers Detection	4
2.5	Generate New Features	5
2.6	Define the Target Metric – Yearly Revenue	5
2.7	Data Codebook.....	6
2	<i>Exploratory Data Analysis Summary</i>	<i>8</i>
3.1	Yearly Revenue	8
3.2	Host Activeness on the Platform	9
3.3	Listings Accommodates & Facilities	13
3.4	Interactive Relationship.....	17
3.5	Geographical Information.....	21
3	<i>Results and In-depth Analysis Using Machine Learning</i>	<i>22</i>
3.1	Preprocess Data	22
3.1.1	Access Collinearity	22
3.1.2	Normalizing & Standardizing	24
3.2	Evaluate & Compare Machine Learning Models.....	24
3.2.1	Models to Evaluate	24
3.2.2	Model Optimization.....	25
4.3	Conclusion & Recommendation	30
4	<i>Conclusion.....</i>	<i>31</i>

1. Introduction

Since its founding, Airbnb has hosted over 60 million people in 34,000 cities across the world and is continuing to grow quickly. Airbnb provides an alternative source of income for people who have otherwise vacant properties and for guests looking for affordable and convenient housing options. With any service, trying to monitor and understand the underlying attributes dynamics of the Airbnb and market trend are very important both for hosts and the entire housing market. As users continue to grow on both the supply and demand side, it becomes crucial to understand and predict how much can a listing earn on the platform.

We seek to analyze over 3,800 listings in the Seattle area in order to better understand how the use of listing attributes such as host activeness, accommodates, location, ratings, and more can be used to accurately predict the optimal earning for the host and those who want to join the platform to become hosts. With better-earning suggestion estimates, Airbnb hosts can improve their management of properties to gain better earning, also the information from the analysis would be insightful for the whole housing market. The objective of this project is to build a model that predicts the earning of a property taking into account listing features and attributes. The end goal is so hosts and the platforms can understand what features of an Airbnb listing are most important to managing and invest for greater yearly.

Answering the question would be a benefit for both potential hosts and any home-sharing platform, that info will be very useful guidance to make the investment as well as regulate the listings market.

2. Data Collection and Wrangling Summary

The detailed data we use [listings.csv](#) (22 February, 2020) was acquired from the Airbnb official website, check the official website for more detailed information. This data set contains 7,544 Observations and 106 features regarding the listing attributes. Each row is an entry for a listing,

and the columns describe the information such as host reaction(host response rate/time), listing properties(property type, room type, bathrooms, bedrooms), and review score. Etc.

We performed data wrangling step with the following data management steps to generate a cleaned dataset for further analysis.

2.1 Drop Features

51 meaningless features are dropped, which includes,

- *32 Meaningless Features* - such as ID, name, constant value, text-related values,
- *16 Redundancy features* - we will keep one features if a multiple same information exist,
- *3 Over-missingness Features* - features with over 50% missing values.

2.2 Re-code the Features

42 features are recoded, which are,

- *Convert the Data Type*- 24 data objects are converted to numeric/categorical/datetime/Boolean data type,
- *Re-categorize the Data* – 17 features are regrouped with less categories,
- *Extract Information from Useable Format* - extract 1 features from list format and recoded as 30 dummy features.

2.3 Deal with Missingness

Dropping cases and imputation were applied,

- 961 (12.7%) records was dropped due to MNAR (Missing Not At Random),
- 2 features were imputed with 0 (13.63% NAs, 5.35% NAs),
- 2 features were imputed by median (9.90% NAs, 0.06% NAs).

2.4 Outliers Detection

Sanity check for all the data range and 0 outliers were dropped.

- Data Range and quantile were calculated for continuous data,

- Percentage for each classes were calculated for categorical data.

2.5 Generate New Features

3 new features were created by the raw features, which are,

- *time_to_first_review*: months a listing takes from join to get the first review,
- *last_review_month*: months a listing get the last review to the time that data was scraped,
- *days_on_airbnb*: days from the first review to the last review.

2.6 Define the Target Metric – Yearly Revenue

This section is a step-by-step demonstration how we define the metric **Yearly Revenue** as the target features.

The [San Francisco Model](#) refers to a method created by Alex Marqusee for the San Francisco Planning Department (Brousseau, 2015) and the Budget and Legislative Analyst's Office (Rodgers, 2015) to quantify the impact of Airbnb in this city. These method given develop ways of estimate business metrics to evaluate local Airbnb listing's booking, occupancy rate and revenue.

We used the San Francisco Model as a guidance for our calculation, below we have listed the features and formulate that leads us to the final business metrics,

- **Days on Airbnb** = $last_{review} - first_{review}$
- **Minimum Booking in Year** = $reviews_per_month * 12$
- **Estimated Booking in Year** = $Minimum\ Booking\ in\ Year / 50\%$
- **Nights Per Year CAP** = $Estimated\ Booking\ in\ Year * minimum_nights_avg_nt$
- **Occupancy Rate** = $Nights\ Per\ Year\ CAP / 365$

- **Yearly Revenue** = *price * Occupancy Rate * 365*

2.7 Data Codebook

The final dataset contains 6,583 observation and 80 features in total. A simplified data codebook is given below,

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 6583 entries, 0 to 7513
Data columns (total 80 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   host_id                                   6583 non-null   int64
1   host_since                               6583 non-null   datetime64[ns]
2   host_response_time                       6583 non-null   object
3   host_response_rate                       6583 non-null   category
4   host_acceptance_rate                    6583 non-null   float64
5   host_is_superhost                       6583 non-null   float64
6   host_identity_verified                   6583 non-null   float64
7   neighbourhood_group_cleansed             6583 non-null   category
8   zipcode                                  6554 non-null   object
9   latitude                                 6583 non-null   float64
10  longitude                                6583 non-null   float64
11  is_location_exact                       6583 non-null   float64
12  property_type                           6583 non-null   category
13  room_type                               6583 non-null   category
14  accommodates                             6583 non-null   int64
15  bathrooms                               6583 non-null   float64
16  bedrooms                               6583 non-null   float64
17  beds                                   6583 non-null   float64
18  price                                   6583 non-null   float64
19  security_deposit                         6583 non-null   float64
20  cleaning_fee                             6583 non-null   float64
21  guests_included                         6583 non-null   int64
22  extra_people                             6583 non-null   float64
23  minimum_nights_avg_ntm                   6583 non-null   float64
24  calendar_updated                         6583 non-null   category
```

25	availability_90	6583 non-null	int64
26	number_of_reviews	6583 non-null	int64
27	number_of_reviews_ltm	6583 non-null	int64
28	first_review	6583 non-null	datetime64[ns]
29	last_review	6583 non-null	datetime64[ns]
30	review_scores_rating	6583 non-null	category
31	review_scores_accuracy	6583 non-null	category
32	review_scores_cleanliness	6583 non-null	category
33	review_scores_checkin	6583 non-null	category
34	review_scores_communication	6583 non-null	category
35	review_scores_location	6583 non-null	category
36	review_scores_value	6583 non-null	category
37	instant_bookable	6583 non-null	float64
38	cancellation_policy	6583 non-null	object
39	require_guest_profile_picture	6583 non-null	float64
40	require_guest_phone_verification	6583 non-null	float64
41	calculated_host_listings_count_entire_homes	6583 non-null	int64
42	calculated_host_listings_count_private_rooms	6583 non-null	int64
43	calculated_host_listings_count_shared_rooms	6583 non-null	int64
44	reviews_per_month	6583 non-null	float64
45	time_to_first_review	6583 non-null	category
46	last_review_month	6583 non-null	category
47	Wifi	6583 non-null	float64
48	Essentials	6583 non-null	float64
49	Heating	6583 non-null	float64
50	Smoke detector	6583 non-null	float64
51	Shampoo	6583 non-null	float64
52	Hangers	6583 non-null	float64
53	Hair dryer	6583 non-null	float64
54	Carbon monoxide detector	6583 non-null	float64
55	Laptop friendly workspace	6583 non-null	float64
56	Iron	6583 non-null	float64
57	TV	6583 non-null	float64
58	Washer	6583 non-null	float64
59	Dryer	6583 non-null	float64
60	Hot water	6583 non-null	float64
61	Fire extinguisher	6583 non-null	float64
62	Dishes and silverware	6583 non-null	float64
63	Refrigerator	6583 non-null	float64
64	Microwave	6583 non-null	float64
65	Coffee maker	6583 non-null	float64
66	Self check-in	6583 non-null	float64
67	Free street parking	6583 non-null	float64
68	Cooking basics	6583 non-null	float64
69	Stove	6583 non-null	float64
70	Bed linens	6583 non-null	float64
71	First aid kit	6583 non-null	float64
72	Oven	6583 non-null	float64
73	Private entrance	6583 non-null	float64
74	Extra pillows and blankets	6583 non-null	float64
75	Free parking on premises	6583 non-null	float64
76	Dishwasher	6583 non-null	float64
77	days_on_airbnb	6583 non-null	float64
78	occupancy_rate	6583 non-null	float64
79	yearly_revenue	6583 non-null	float64

dtypes: category(14), datetime64[ns](3), float64(51), int64(9), object(3)
memory usage: 3.5+ MB


More details and source codes on Data Wrangling are available at
[Capstone_Project_Data_Wrangling](#).

2 Exploratory Data Analysis Summary

In this section, we perform the Exploratory Data Analysis on the cleaned data generated from the Data Wrangling step. We will go through most of the features in the data to explore their relationship with the Yearly Revenue. And the EDA is organized based on different topics on the related information, features regarding the same topic information will be analysis together.


3.1 Yearly Revenue

The Yearly Revenue is the target business metric we want to evaluate for all the listing population from our data, the definition was given in 2.7. Let's evaluate the distribution and statistics summary for it.



Yearly Revenue Statistics

Mean	42788.59
Std	92133.01
Min	96.00
25%	6460.56
50%	14968.80
75%	35164.80
Max	1659240.00




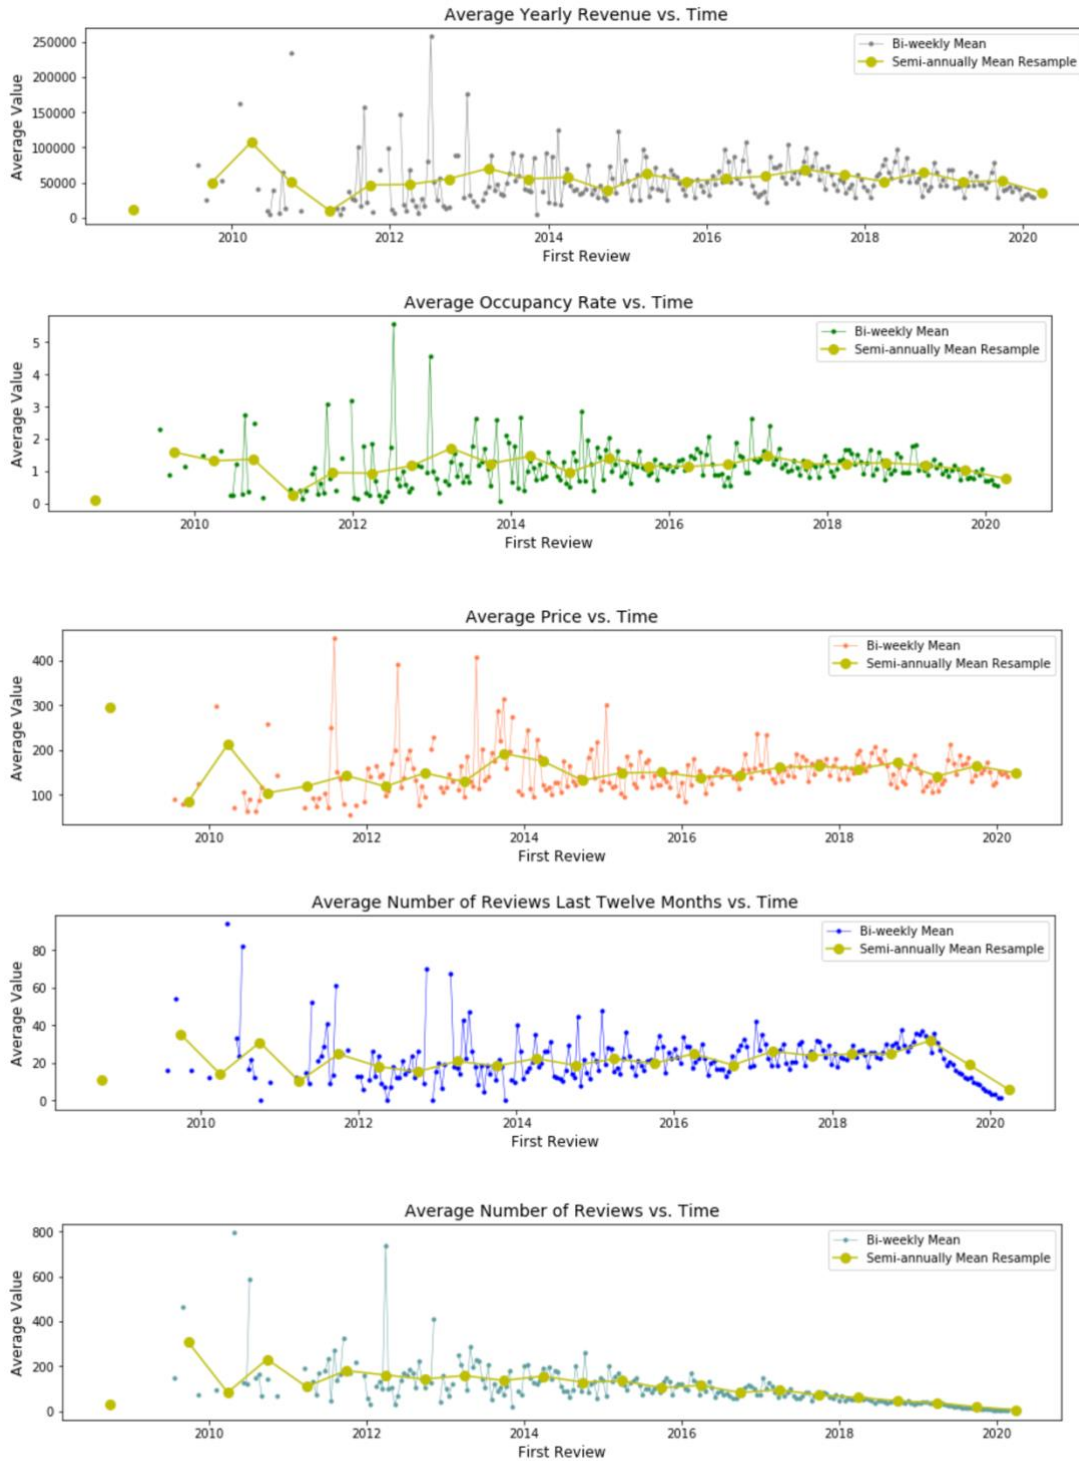
From the summary statistics above, we can see that a listing can have no revenue at all among the year, also a listing can gain a very high revenue up to \$1,659,240. However, the distribution of the yearly revenue is highly right-skewed, we can see that most of the data are centered within \$100,000, and 99% of our data lies within \$500,000. The median is about \$14,968 and the average revenue goes up to \$42,788. The density plot is given alongside. Putting those numbers into perspective, the 2019 Airbnb official sites reported that On average globally, Experience hosts earn \$10,000 a year, our data reveals a info that indicating a healthy premium for actively managed listings for the Seattle Airbnb Market.

3.2 Host Activeness on the Platform

In this section, we will exam the host activeness from different perspectives, which includes the global active on the platform, the ability to attract the guest and stay on the market, and it's longevity on the platform. And also explore the their relationship with the yearly revenue.

Global Business Trend

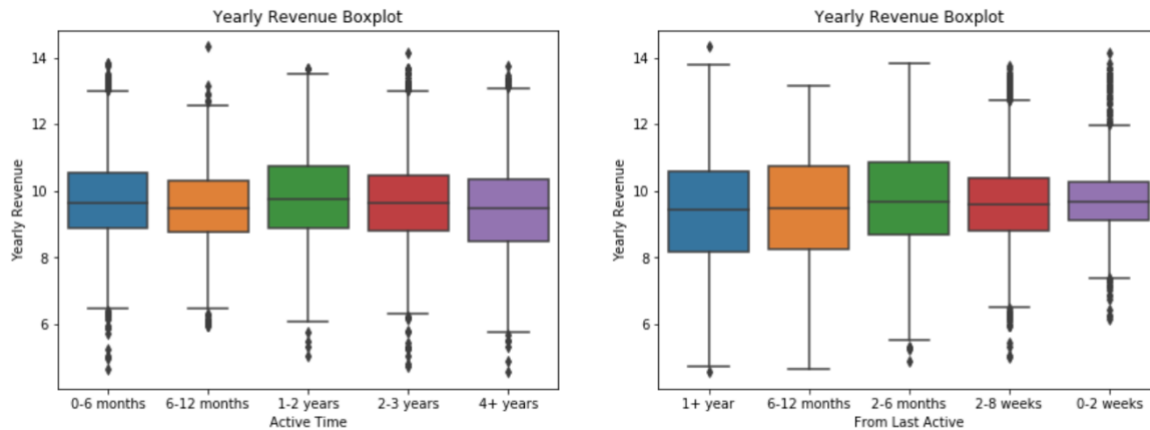




The first plot indicates that the business is keep boosting all these years – more host join Airbnb and the number of host get its first review is increasing. In general, during the decades of 2010 to 2020, there are not tremendous increase among all the business metrics. The yearly revenue, price, and occupancy rate look flat with a slightly increasing tendency. Noticed that the both

the average review number and that feature in the last twelve month have a slightly decrease trend could cause by delay of getting the review data coming in.

Power of Attraction & Staying






Statistical summary table

active_time				from_last_active			
	Count	Mean	Median		Count	Mean	Median
0-6 months	1,603	\$47,808	\$15,585	0-2 weeks	1,655	\$33,775	\$16,236
6-12 months	630	\$34,332	\$12,994	2-8 weeks	1,987	\$43,877	\$14,515
1-2 years	981	\$51,735	\$17,280	2-6 months	1,463	\$54,541	\$16,245
2-3 years	1,743	\$43,588	\$15,552	6-12 months	480	\$41,594	\$13,212
4+ years	1,626	\$34,860	\$13,058	1+ years	998	\$38,912	\$12,391


The created feature active_time is the months values from join to get the first review, which indicated power of attraction with shorter time. It looks like the among the groups of different months length to get the first review, the yearly revenue doesn't change much. The statistics summary indicates that the 1-2 years have the highest mean (\$51,735) and median (\$17,280). To embody the staying power, we created feature from_last_active, which is the length of last review to the data was scrapped time. The plots indicates that as the better staying power (get the last review as recent as it can), the higher the average yearly revenue.

Host Listing Counts



The Host Listing Count features are specified for entire homes/private rooms/shared rooms, the average yearly revenue varies and show no specific patterns among different counts.

Days on Airbnb




The features days_on_airbnb is calculated length from first review to last review, which potentially indicates the activeness length on the platform. Original pearson correlation is weak

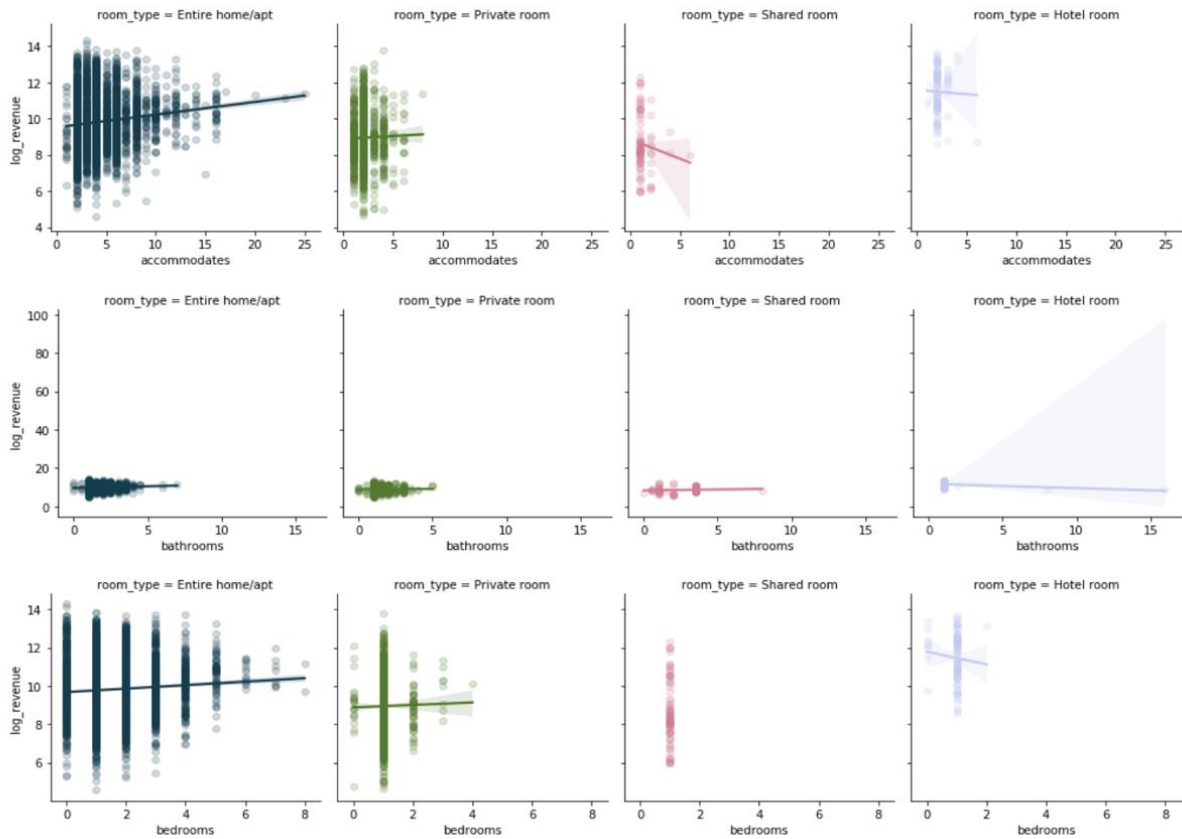
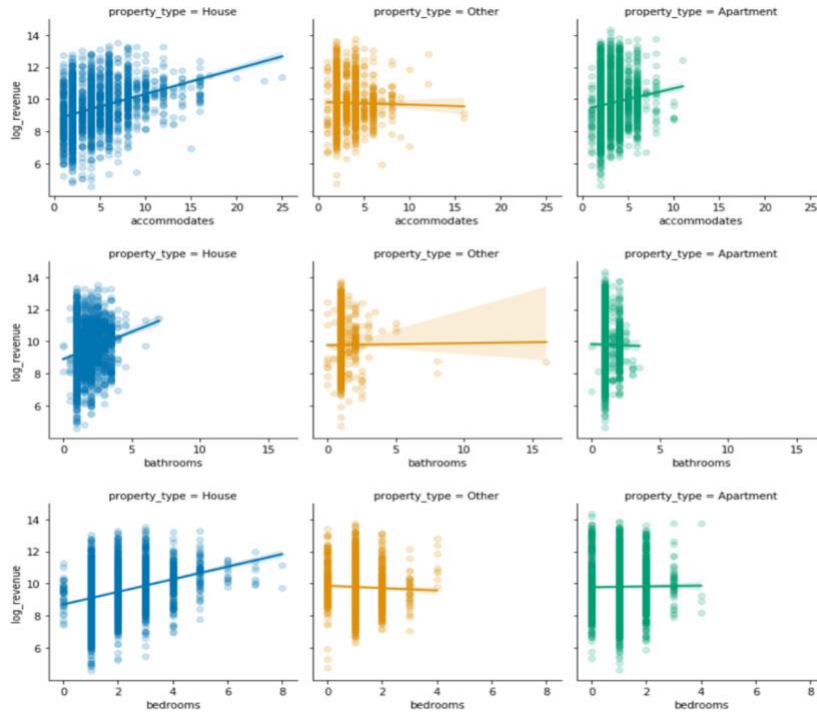
0.035 between the yearly revenue and days on Airbnb. The KDE plot above show the log transformed revenue and correlation. We can see that most intense data are within 0-1000 days, which is about 3 years, there's no indication that longer days is associated with higher revenue.

3.3 Listings Accommodates & Facilities

Bedrooms & Bathrooms & Accommodates




The correlation matrix above shows negligible correlation between the revenue and bedrooms/bathrooms, let's take a further examination this accommodates/bedrooms/bathrooms features within two features - property type and room type.



The above detailed plots the relationship between yearly revenue with each class in property type and room type. Only under the property type is house and apartment, there's slightly positive relationship indicates that as the number of accommodates increases, the yearly revenue increases. And the same for revenue and accommodates within property type is apartment.

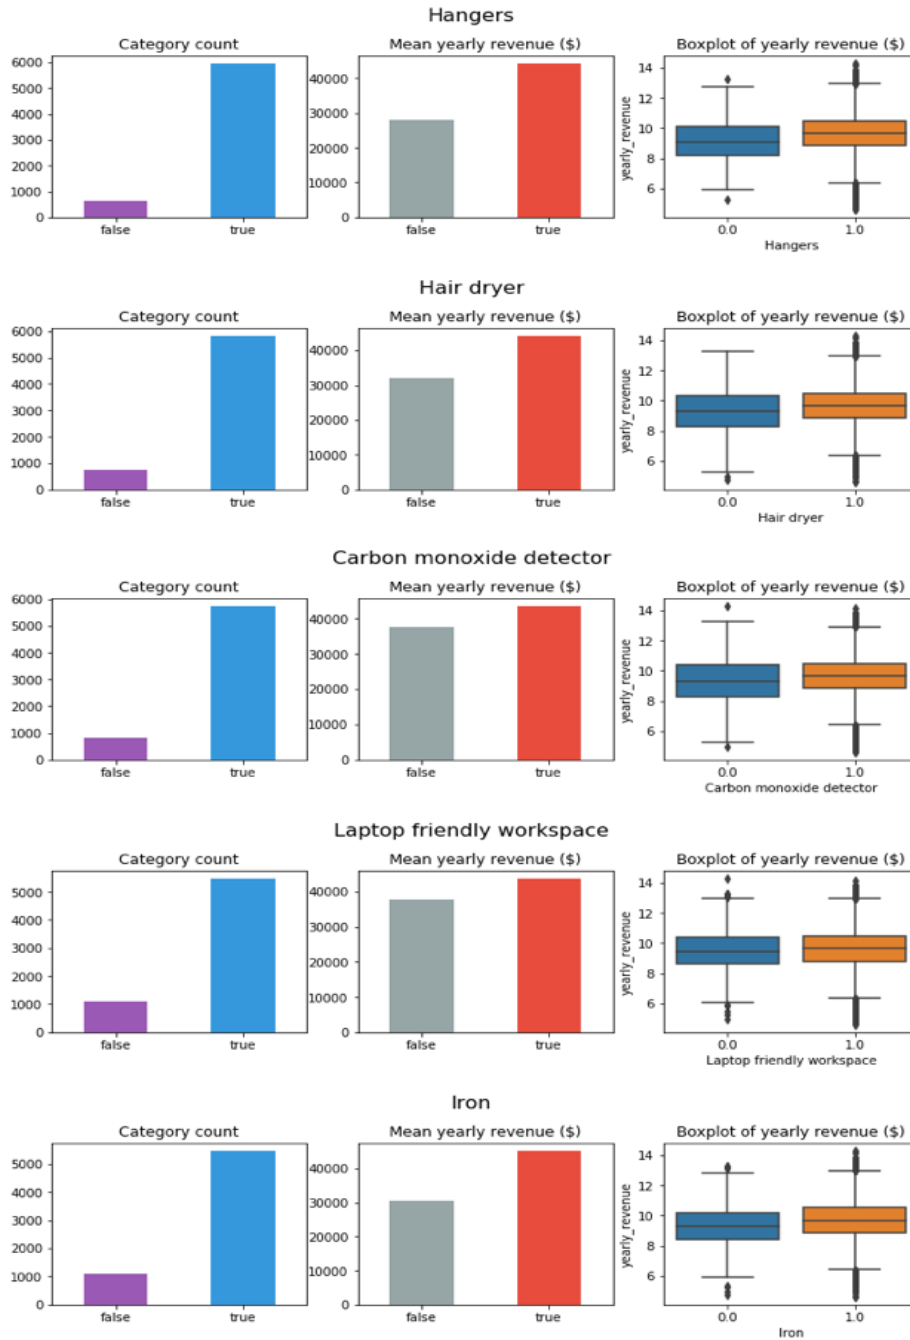
Clean Fee & Guest Included & Availability



The above features are examined with log transformed yearly revenue. There is minor positive relationship between the yearly revenue and features of cleaning fee and guests included.

Amenities

We have 30 single amenity features to check, they are essential facilities like WIFI, TV, Hairdryer, etc. We combine the plots of the counts with and without the single amenities, the mean revenue for each class along with the boxplot, we display the 5 out of 10 single amenities here, check the Jupyter notebook for more details.



For better information, we implement bootstrap hypothesis testing across the 30 amenities. Univariate test reveals statistical significance for five amenities - **Fire extinguisher, Free street parking, First aid kit, private entrance and Free parking on premises**, which indicates with those amenities provided, the yearly revenue are statistically different. Further examination brings in more interesting results - the mean yearly revenue are lower with those amenities.



Bootstrap hypothesis test for difference of means yearly revenue:



Hypothesis	Test of Wifi	p-value = 0.4641
Hypothesis	Test of Essentials	p-value = 0.9673
Hypothesis	Test of Heating	p-value = 0.4547
Hypothesis	Test of Smoke detector	p-value = 0.9892
Hypothesis	Test of Shampoo	p-value = 0.9614
Hypothesis	Test of Hangers	p-value = 1.0
Hypothesis	Test of Hair dryer	p-value = 1.0
Hypothesis	Test of Carbon monoxide detector	p-value = 0.9708
Hypothesis	Test of Laptop friendly workspace	p-value = 0.9898
Hypothesis	Test of Iron	p-value = 1.0
Hypothesis	Test of TV	p-value = 1.0
Hypothesis	Test of Washer	p-value = 1.0
Hypothesis	Test of Dryer	p-value = 1.0
Hypothesis	Test of Hot water	p-value = 0.7468
Hypothesis	Test of Fire extinguisher	p-value = 0.0
Hypothesis	Test of Dishes and silverware	p-value = 0.2256
Hypothesis	Test of Refrigerator	p-value = 0.0344
Hypothesis	Test of Microwave	p-value = 0.0519
Hypothesis	Test of Coffee maker	p-value = 0.4128
Hypothesis	Test of Self check-in	p-value = 1.0
Hypothesis	Test of Free street parking	p-value = 0.0
Hypothesis	Test of Cooking basics	p-value = 0.9986
Hypothesis	Test of Stove	p-value = 0.9977
Hypothesis	Test of Bed linens	p-value = 0.9245
Hypothesis	Test of First aid kit	p-value = 0.0
Hypothesis	Test of Oven	p-value = 0.9998
Hypothesis	Test of Private entrance	p-value = 0.0
Hypothesis	Test of Extra pillows and blankets	p-value = 0.798
Hypothesis	Test of Free parking on premises	p-value = 0.0
Hypothesis	Test of Dishwasher	p-value = 0.9967

3.4 Interactive Relationship

The section we group the features into two perspectives and the exploration analysis will be performed from two sides.

From Host's Side







The more active response rate and time are not appear to have a better revenue than we expected. Neither are the features as cancellation policy and calendar updated time.

We explore the five features – *host_is_superhost*, *is_location_exact*, *instant-bookable*, *require_guest_profile_picture* and *require_guest_phone_verification* below. The visualization and bootstrap hypothesis testing were applied for the

the It looks like only require guest profile picture and require quest phone verification are associated with lower yearly revenue – maybe the guests prefer no verification and skip these listing when making a choice.

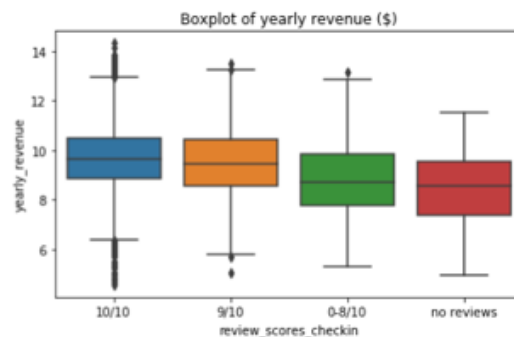

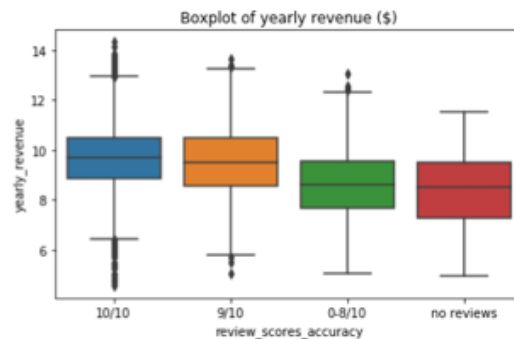
Bootstrap hypothesis test for difference of means yearly revenue:

Hypothesis Test of <i>host_is_superhost</i>	p-value = 1.0
Hypothesis Test of <i>is_location_exact</i>	p-value = 0.4659
Hypothesis Test of <i>instant_bookable</i>	p-value = 1.0
Hypothesis Test of <i>require_guest_profile_picture</i>	p-value = 0.0
Hypothesis Test of <i>require_quest_phone_verification</i>	p-value = 0.0






From Guest's Side

The features we can get info from the guest's side are their rating scores. There are 7 review scores regarding accuracy, cleanliness, checking, communication, location and value. As we expected – the higher score group earn more revenue in general, and the yearly revenue trend almost monotonic decreasing along with the review score go down. Noted that 3 features are given here, check the Jupyter notebook for more details.



3.5 Geographical Information

We mapped the median revenue that group by the geographical information (latitude, longitude). It's obvious that more close to the central part, the revenue is getting higher. When zoom in the plot, it's more clear some region such as Queen Anne, and Capitol Hill.



We group the data by the neighborhood classes and use bar plot to display the average revenue for each area, the trends align with we get from the geographical data, the top regions are Cascade, Capitol Hill, Queen Anne, Downtown, and etc.

3 Results and In-depth Analysis Using Machine Learning

This section consisted of two parts, we will start from the **Data Preprocessing**, in which we will access the collinearity to drop the highly correlated features, and also apply the normalizing and standardizing on the dataset. After that, we will apply the machine learning algorithms in the **Modeling** part with the processed data.

3.1 Preprocess Data

3.1.1 Access Collinearity

We use the heatmap as a tool to access the correlation between all the features we have.



Figure. Heatmap on features before dropping

The heatmap show's before processing, the correlation range is (-0.75, 0.75). Some features are highly correlated,

- *Accommodates and bathrooms, bedrooms, bed, guest_included,*
- *require_guest_phone_verification and require_guest_profile_picture,*
- *Dryer and Washer, Dishes and silverware and Refrigerator etc.,*
- All the '*No Reviews*' features, all the '*review_score_9/10*'

For all of the correlated features, we will only keep one from them. The heatmap for the reduced data are given below, the correlation range go down to (-0.6, 0.6).



Figure. Heatmap on features after dropping

3.1.2 Normalizing & Standardizing

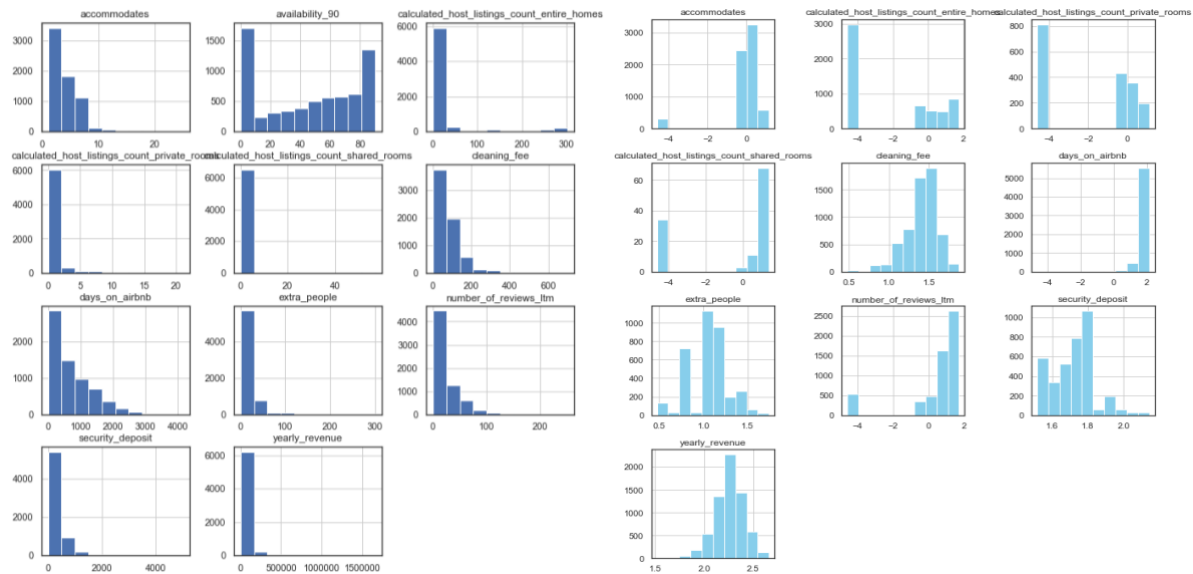


Figure. Before vs. after normalization

Then for all the continuous features, we applied the log transform. We still see after that not every single features approached the normal distribution perfectly, but all of them improved a lot.

Finally, we splitted the data into train (80%) and test (20%) dataset, and the standard scaler is applied to the train data, the train scaler the applied to test data to prevent the leaking problem.


3.2 Evaluate & Compare Machine Learning Models

3.2.1 Models to Evaluate

We will compare five different machine learning models using the Scikit-Learn library:

- Linear Regression
- Support Vector Machine Regression
- Random Forest Regression
- Gradient Boosting Regression
- K-Nearest Neighbors Regression

To compare the models, we are going to be mostly using the Scikit-Learn defaults for the model hyperparameters. And in this section, we will focus on these models' baseline performance on our dataset generally instead of optimizing the model with a determined model to use. Then we can select the best performing model for further optimization using hyperparameter tuning.



We use the testing Mean Absolute Error (MAE) and Mean Standard Error (MSE). The Base model comparison shows Random Forest, Gradient Boost and Linear Regression have the best performance.

3.2.2 Model Optimization

3.2.2.1 Hyperparameter Tuning

In machine learning prediction, optimizing a model will find us the best set of hyperparameters for to achieve the best performance. we will start with the best base model random forest first.

In our case of a random forest, hyperparameters include,

- **n_estimators** - number of trees in the foreset
- **max_features** - max number of features considered for splitting a node
- **max_depth** - max number of levels in each decision tree
- **min_samples_split** - min number of data points placed in a node before the node is split
- **min_samples_leaf** - min number of data points allowed in a leaf node
- **bootstrap** - method for sampling data points (with or without replacement)

The best hyperparameters are usually impossible to determine before training, which means the best practice to determine the optimal settings is to try many different combinations to evaluate the performance of each model. However, evaluating each model only on the training set can cause overfitting - which means our model will score very well on the training whereas it will not be able to generalize to new data as good as we expected.

Hyperparameter Tuning with Random Search within Cross Validation will generate more reliable results for our reference.

In our case, we will use Random Search with cross validation to explore where the model can have relatively good performance, once we narrow down the options and then use the grid search with a more limited range of options for the best model.

The K-Fold Cross Validation technique dividing the training data into K folds, and then going through an iterative process on each K-1 fold data set and evaluate performance on the kth fold. Average error on each of the K iterations as the final performance will be computed as the given results.

The next section will demonstrate how we use **random search** to narrow down the scope and use **grid search** to find our best model, with the combination of cross validation.

Optimization Parameters

Optimization	Search Range	Best parameters	Cross - Validation
Random Search	n_estimators: (200, 2000) by 10	n_estimators: 400	3 fold
	max_features: [auto, sqrt]	max_features: sqrt	
	max_depth: (10, 110)	max_depth: None	
	min_samples_split: [2, 5, 10]	min_samples_split: 2	
	min_samples_leaf: [1, 2, 4]	min_samples_leaf: 1	
	bootstrap: [True, False]	bootstrap: False	
Grid Search	n_estimators: (200, 700) by 50	n_estimators: 400	3 fold
	max_features: sqrt	max_features: sqrt	
	max_depth: None	max_depth: None	

min_samples_split: 2

min_samples_split: 2

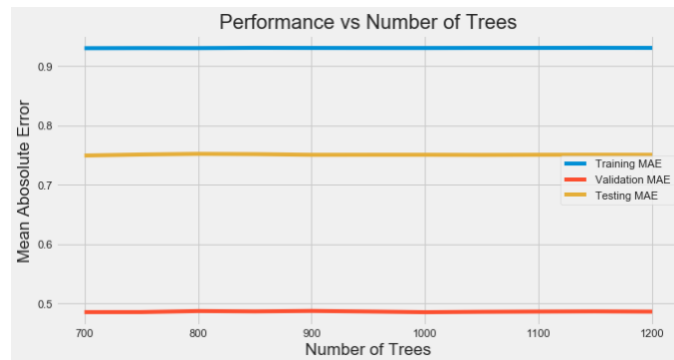
min_samples_leaf: 1

min_samples_leaf: 1

bootstrap: [True, False]

bootstrap: [True, False]


To proceed forward the grid search best model, we also plot the training error, validation error and testing error to see if the model have overfitting issue.



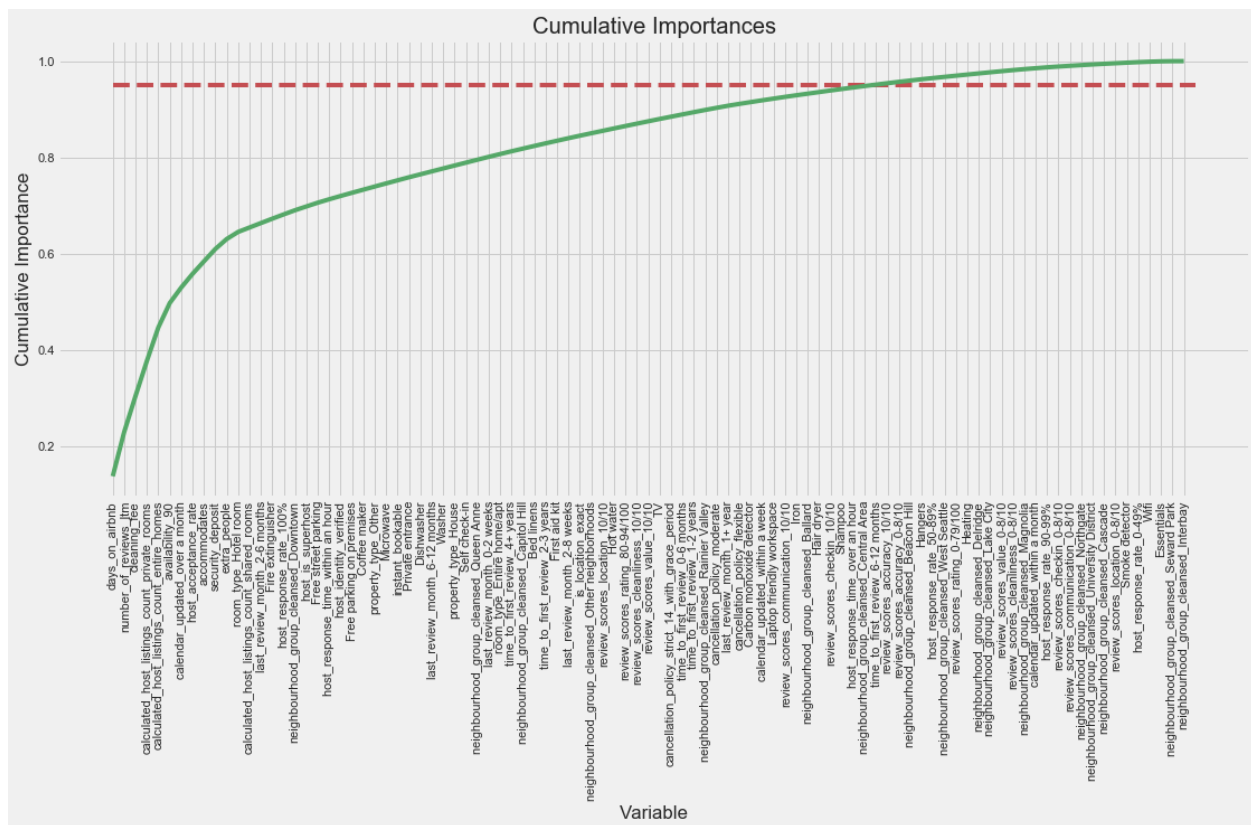
Within cross validation, we use 2 batch of data to train the model and 1 data set for validation. So the cross validation results are good – the validation error are lower than the training error. However, when the model is applied to the testing data, the model is involved with **overfitting** issue.

4.2.2.2 Feature Reduction

To better reduce the overfitting issue – we think of the curse of dimensionality – our model have 93 features! It's time to consider the feature reduction for less complex model.



We plot the top 50 in our model, we can see down to the very bottom, some features have minimum vote for the model.



Further, we plot the accumulative importance by descending sorted features, and we added a 95% cutoff line, all the features will be used in the new model with the best parameter in precious session as the re-fitting model.


After refitting, the model contains 68 features, and the performance is as good as the model with 96 features.

Model Optimization Results

Model	MAE	MSE	Accuracy
Base Model	0.75	1.04	90.27%
Random Search best model	0.65	0.98	92.07%
Grid Search best model	0.76	0.98	92.01%
Feature Reduced model	0.75	1.04	92.06%


Since we already know that from the first beginning that the linear regression has a relatively good performance, we want to keep the model as simple as it can be, so we will use the reduced model to fit the linear regression.

The fitted linear model has MAE as 0.91, MSE as 1.38 and Test Accuracy as 90.36, which is very close to our featured reduced random forest model. Let's check the model assumption.



In general our model have good performance, however, it have underprediction for the peak revenue value and overprediction for the tail in the two sides.

4.3 Conclusion & Recommendation



Compared to the best random forest model with an accuracy of 89.38%, the ridge regression would be a better choice for predicting the revenue.

Model	Feature Numbers	Prediction Accuracy	Accuracy- CV	Overfitting
Ridge - Reduced Features	58	96.21%	89.6%	No
Random Forest - Best Random Search	102	89.64%	95.84%	Yes
Random Forest - Reduced Features	58	89.64%	-	-
Random Forest - Best Grid Search	102	89.61%	95.9%	Yes
Random Forest - Baseline	102	88.68%	-	-

Based on the results above, the linear would be our final model, which has,

- best Prediction accuracy
- no overfitting
- a relatively simple algorithm for a further system implementing

4 Conclusion